



SC Asia 2018 (28/3/2018)

MACC: An OpenACC Transpiler for Automatic Multi-GPU Use

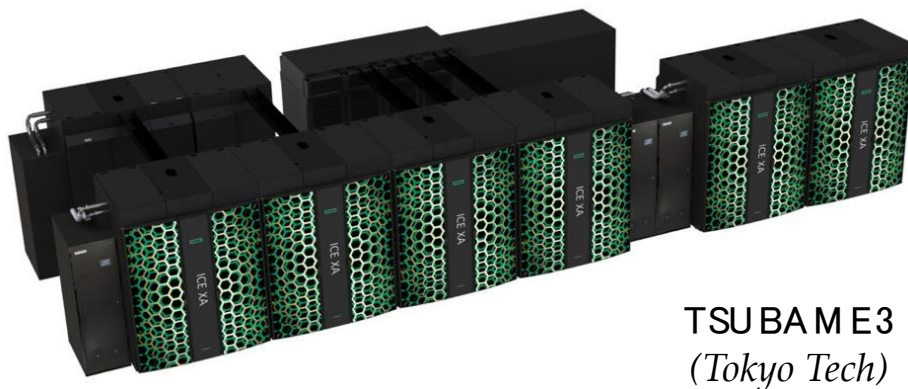
Kazuaki Matsumura^{1,2}, Mitsuhsa Sato^{3,4}, Taisuke Boku⁴

Artur Podobas¹, Satoshi Matsuoka^{1,2}

¹ Tokyo Institute of Technology ² AIST-Tokyo Tech RWBC-OIL ³ RIKEN AICS ⁴ University of Tsukuba

Introduction

- Modern supercomputers are heterogeneously designed
 - Execution cooperatively with Accelerators (e.g. GPU, Intel MIC, FPGA)
- Several state-of-the-art supercomputers contain multiple GPUs (multi-GPU)



- Using accelerators incurs additional programming cost
 - Through primitives (CUDA, OpenCL) or abstract models (DSL, Directive-based)

What is OpenACC ?

- Directive-based programming models complement naive code by putting directives
- OpenACC is a directive-based programming model developed by OpenACC organization
 - Supported by PGI Compiler (*NVIDIA*) and Cray Compiler, and experimentally by GCC
- Realizes accelerator execution by inserting directives into original C / Fortran source-code,



Example: N-Body Computation

■ Communication Code
■ GPU Invocation Code

CPU Code

```

cudaMalloc(&dev_m, N * sizeof(float));
cudaMalloc(&dev_p, N * sizeof(float3));
cudaMalloc(&dev_v, N * sizeof(float3));
cudaMemcpy(dev_m, m, N * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(dev_p, p, N * sizeof(float3), cudaMemcpyHostToDevice);
cudaMemcpy(dev_v, v, N * sizeof(float3), cudaMemcpyHostToDevice);

for (int t = 0; t < TIME_STEP; t++) {
    kernel1<<<block_num, thread_num>>>(dev_m, dev_p, dev_v);
    kernel2<<<block_num, thread_num>>>(dev_m, dev_p, dev_v);
}
    
```

GPU Code

```

__global__ void kernel2(float *m, float3 *p, float3 *v){
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    int offset = tid * THREAD_SIZE;
    for (int j = 0; j < THREAD_SIZE; j++) {
        int i = offset + j;
        p[i].x += v[i].x * DT;
        p[i].y += v[i].y * DT;
        p[i].z += v[i].z * DT;
    }
}
    
```



CPU + GPU Code

```

#pragma acc data copyin (p_x[N], p_y[N], p_z[N], m[N])
#pragma acc data copyout (v_x[N], v_y[N], v_z[N])
for (int t = 0; t < TIME_STEP; t++) {
    #pragma acc parallel loop independent
        for (int i = 0; i < N; i++) { /* ... */ }

    #pragma acc parallel loop independent
        for (int i = 0; i < N; i++) {
            p_x[i] += v_x[i] * DT;
            p_y[i] += v_y[i] * DT;
            p_z[i] += v_z[i] * DT;
        }
}
    
```

OpenACC
 Directives for Accelerators

Motivating example: Multi-GPU with OpenACC

Single-GPU Code

```
#pragma acc data\
    copyout(x[0:N]) present(y)
#pragma acc kernels
for (int i = 0; i < N; i++)
    x[i] = y[i] * y[i];
```

- + Concurrent Execution
- + Data Transfer
- + Loop Division

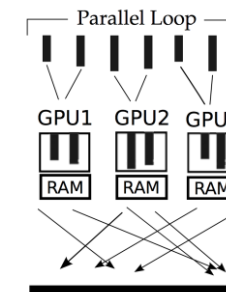
```
numgpus = acc_get_num_devices(DEVICE_TYPE);
#pragma omp parallel num_threads(numgpus)
{
    int tnum = omp_get_thread_num();
    int sz = N / numgpus;
    int lb = sz * tnum; int ub = lb + sz;
    acc_set_device_num(tnum, DEVICE_TYPE);
    #pragma acc data copyout(x[lb:sz]) present(y)
    #pragma acc kernels
    for (int i = lb; i < ub; i++)
        x[i] = y[i] * y[i];
}
```

Multi-GPU Code

- Manual efforts break out of the abstraction
- We propose an automation method which requires no modification of original OpenACC source-code

Related Work

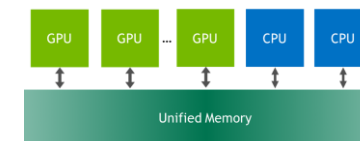
“Integrating multi-GPU execution in an OpenACC compiler” (ICPP’13)



Quoted figure: 1)

- Keeps the coherence of array chunks: *Bulk Synchronous Parallel model*
- Programmer has to provide the chunk size and additional annotations to optimize communications

Unified Memory of NVIDIA GPUs



Quoted figure: 2)

- Keeps all data coherent without user intervention → The problem is overheads.

“Automatic data allocation and buffer management for multi-GPU machines” (TACO’13)

- Leverages the Polyhedral Model (a strict affine model) to detect fine dependencies

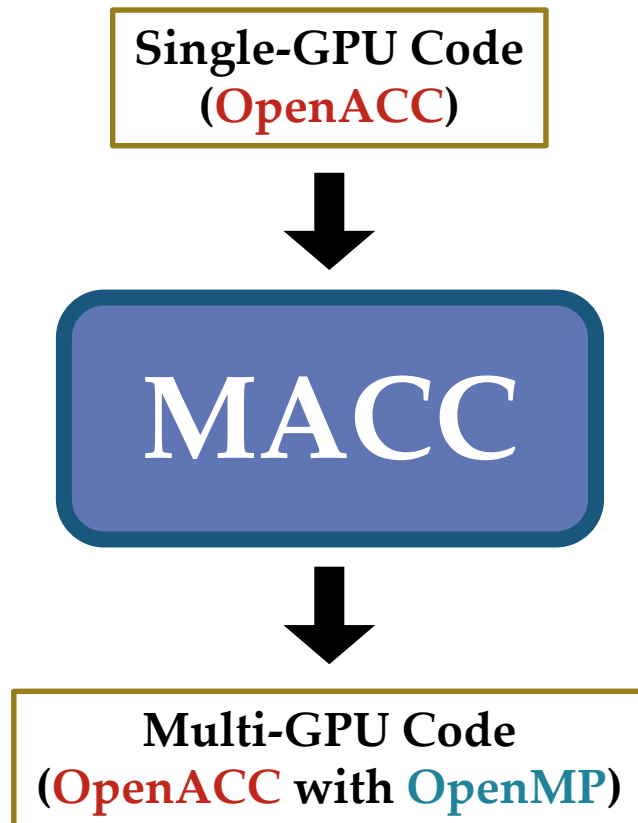
$$\text{X} A[B[s]] = 1;$$

1) Komoda Toshiya, Shinobu Miwa, Hiroshi Nakamura, and Naoya Maruyama. Integrating multi-GPU execution in an OpenACC compiler. In the 42nd International Conference on Parallel Processing (ICPP), 2013.

2) <https://devblogs.nvidia.com/unified-memory-cuda-beginners/>

3) Thejas Ramashekar, and Uday Bondhugula. Automatic data allocation and buffer management for multi-GPU machines. In ACM Transactions on Architecture and Code Optimization (TACO), Vol. 10, No. 4, Article 60, 2013.

Proposal: A Transpiler for Automatic Multi-GPU Use



- We propose a transpiler (source-to-source compiler) named MACC
- The output code can exploit multi-GPU without any manual effort, keeping original semantics and portability: **OpenACC** + **OpenMP**
(GPU Parallel) (CPU Parallel)
- This proposal has a generality to support:
 - Multiple accelerators (not only GPUs)
 - Other directives for accelerators (e.g. **OpenMP**'s target directive)

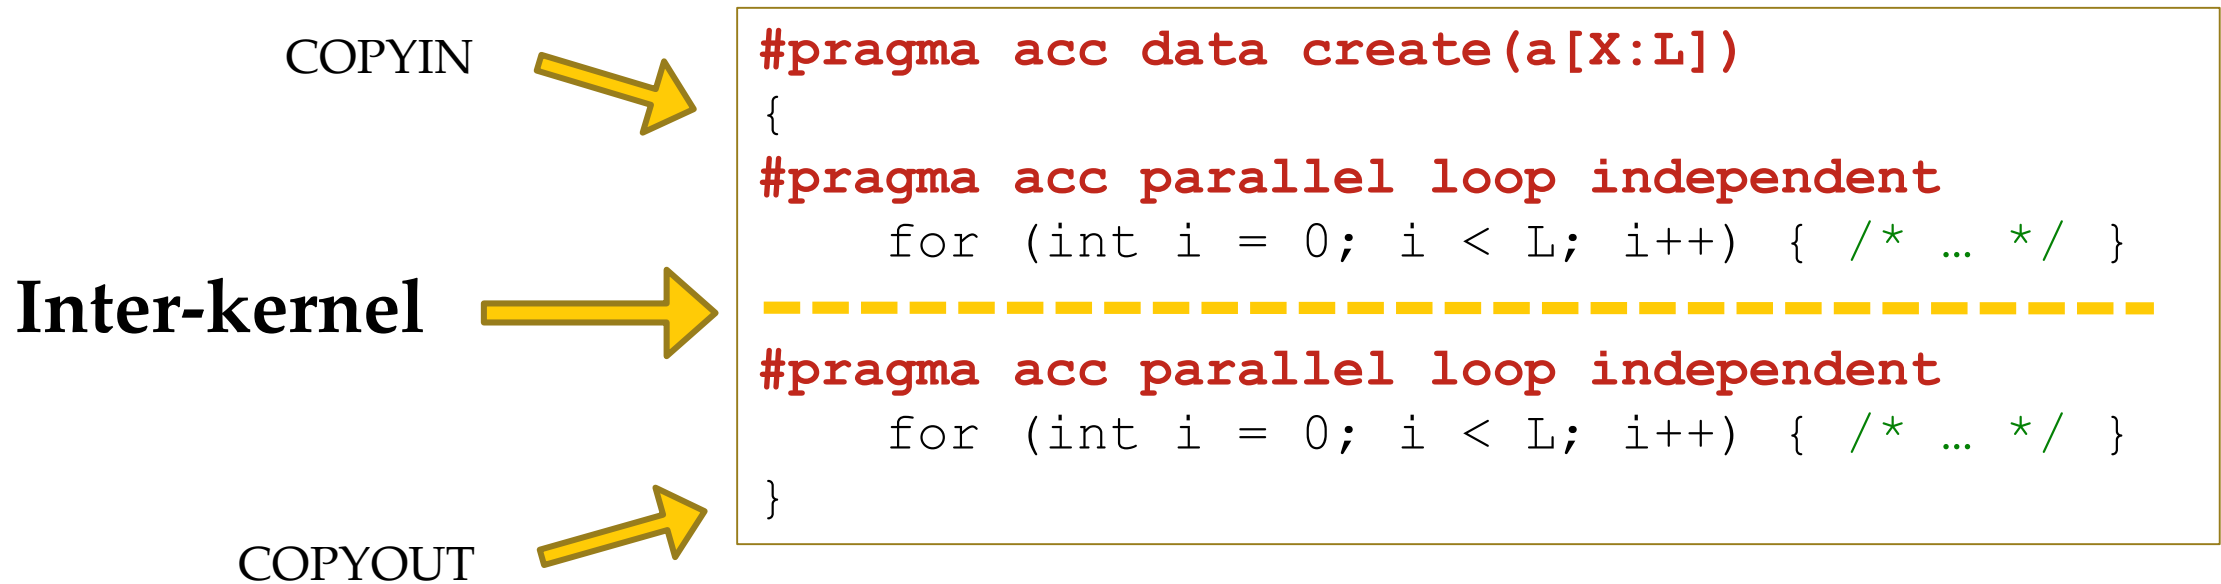
Strategy: How to parallelize kernels over multi-GPU?

Assumption 1. OpenACC primarily targets loop-level parallelism of the outermost loop

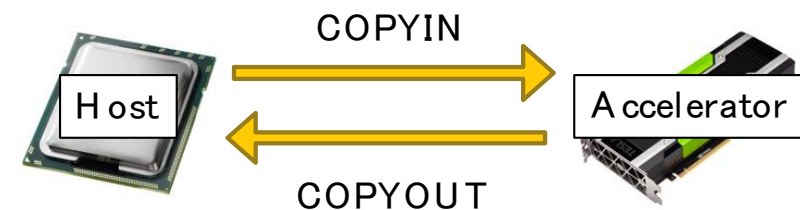
Assumption 2. Well-written OpenACC programs have continuous accesses

→ Equally divide the outermost loop for each GPU. If not dividable, execute on single-GPU.

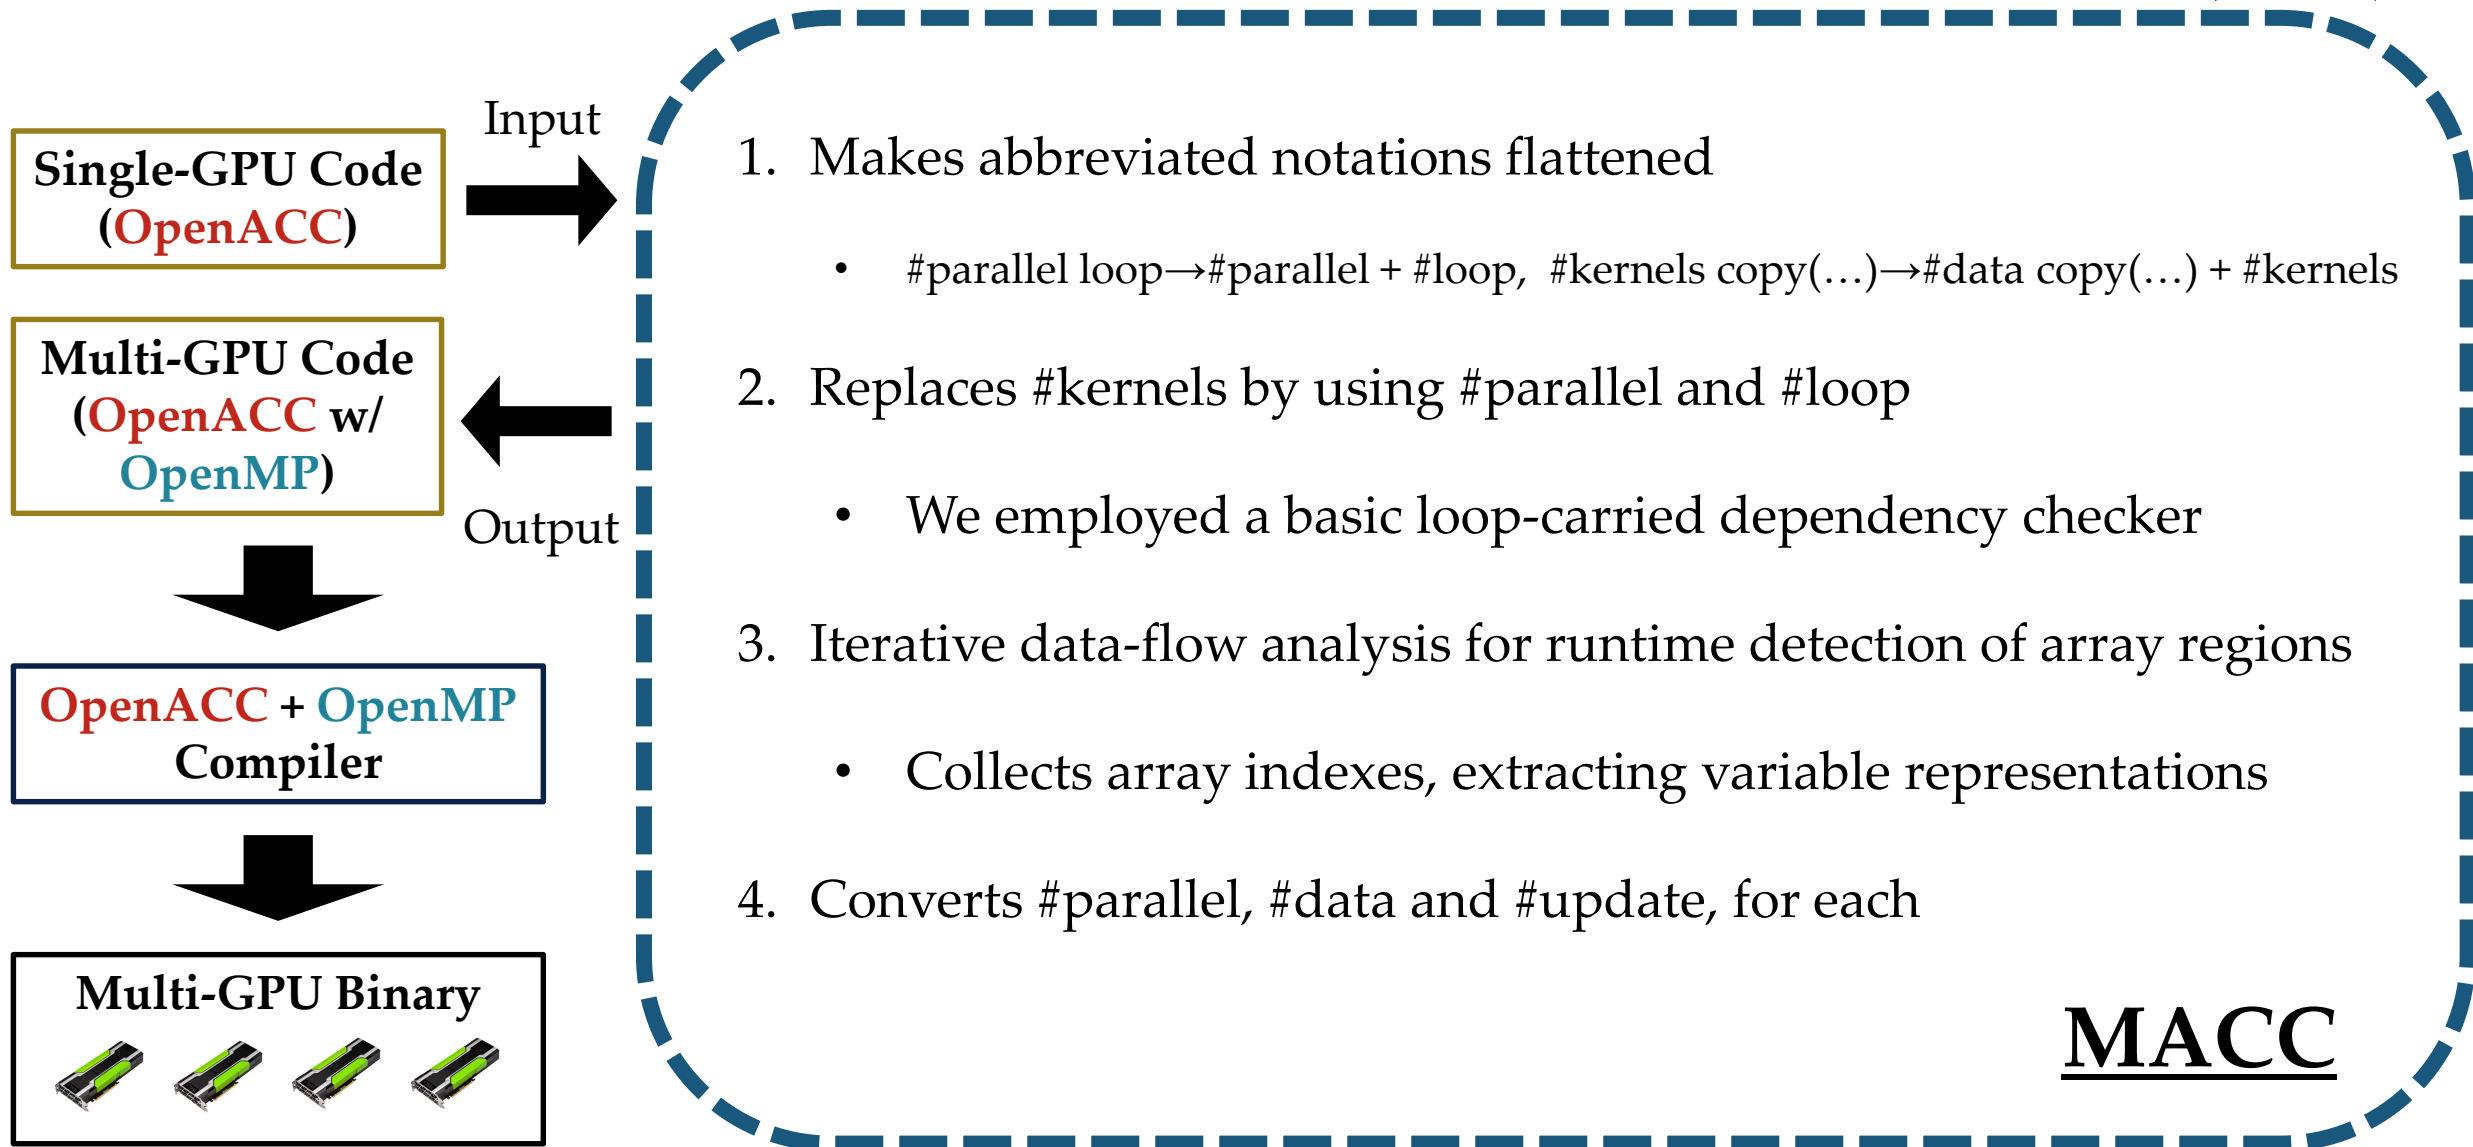
Challenge: Inter-kernel communication



- Data-dependency among GPUs must be solved at the inter-kernel communication
- We utilize upper/lower-bounds of array accesses to generate communications automatically



MACC: Entire Process



```
#pragma acc parallel
{ /* ... */ }
```

```
if (/* sections are changed */)
{ /* recalculate sections */ }
#pragma omp parallel num_threads(NUMGPUS)
{
    int tnum = omp_get_thread_num();
    set_gpu_num(tnum);
    set_data_section(/* ... */);
#pragma omp barrier
#pragma acc parallel
    { /* Splitted Loop */ }
}
```

```
#pragma acc data\
    copy(x[0:N])
{ /* ... */ }
```

```
#pragma omp parallel num_threads(NUMGPUS)
{
    copyin_routine(omp_get_thread_num(), x, 0, N);
}
{ /* ... */ }
#pragma omp parallel num_threads(NUMGPUS)
{
    copyout_routine(omp_get_thread_num(), x);
}
```

```
#pragma acc update\
    host(x[a:b])
```

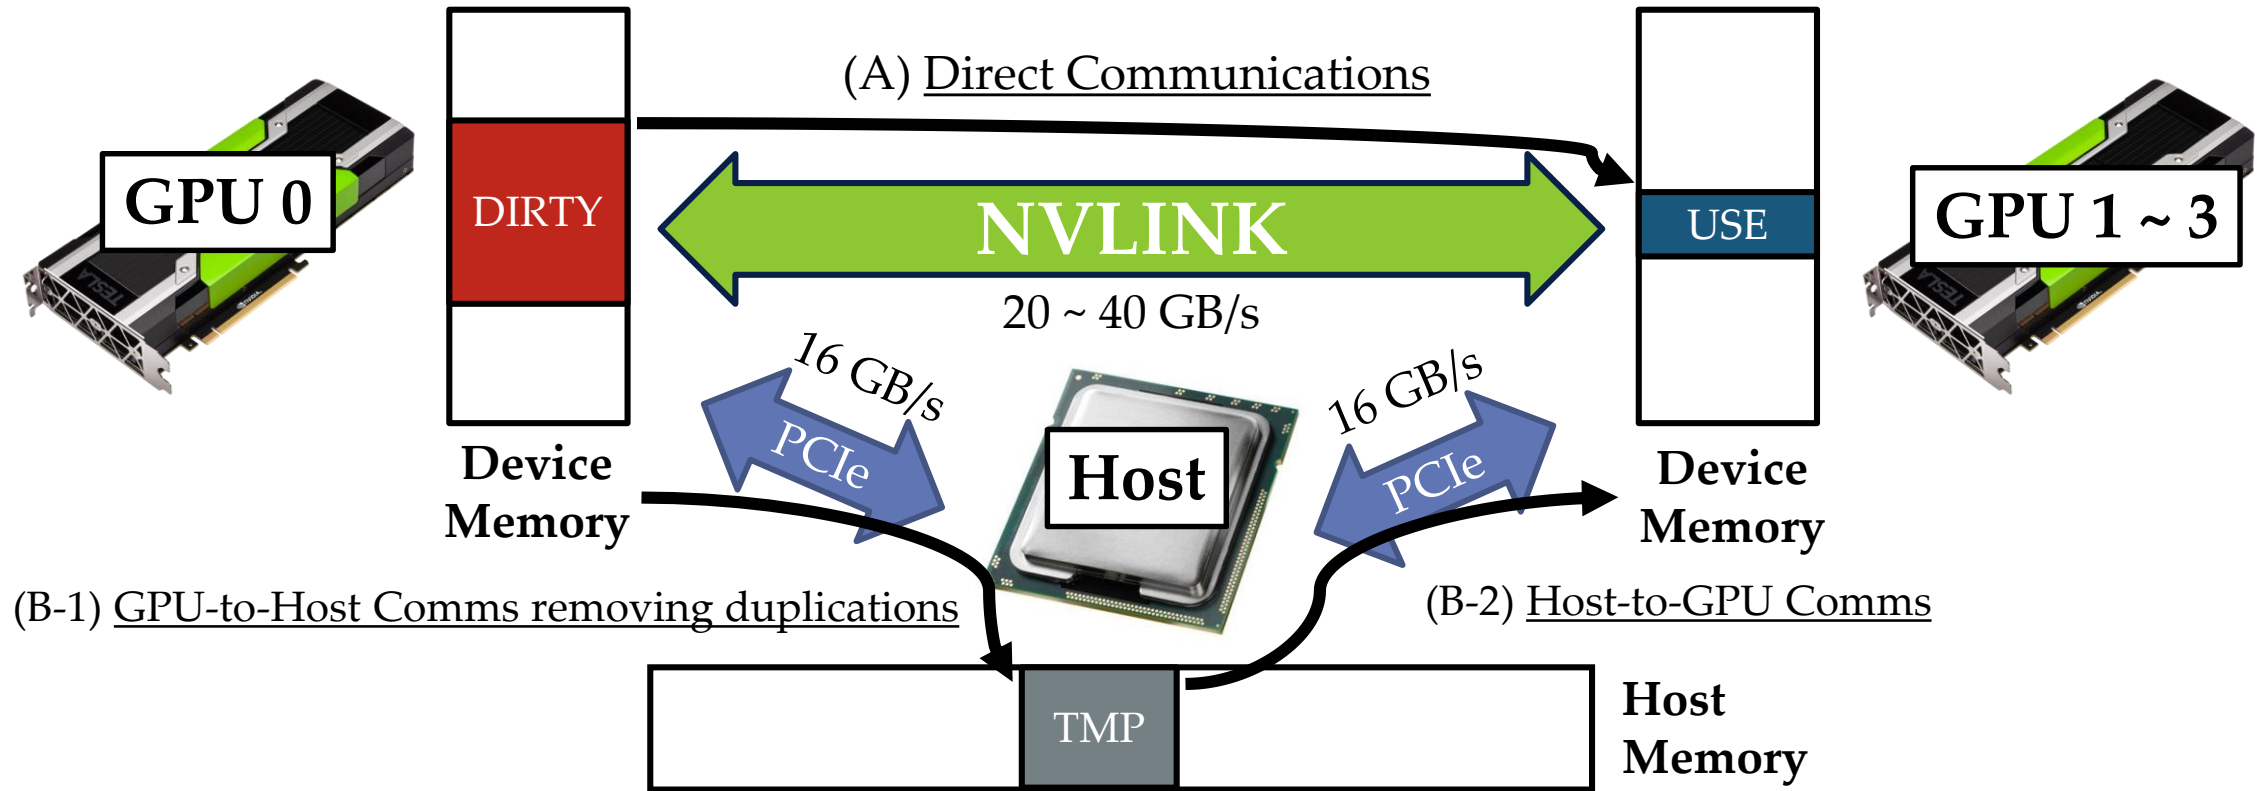
```
#pragma omp parallel num_threads(NUMGPUS)
{
    int tid = omp_get_thread_num();
    update_host_routine(tid, x, a, b);
}
```

- Convert to use MACC's communication routines
 - Which call OpenACC's routines
- **OpenMP** is for parallelizing multi-GPU execution
 - No restriction to be replaced by POSIX threads or **OpenACC's** Async
 - But it benefits to:
 - Code generation (easy)
 - Better performance (than Async)
 - Multi-core execution (supported by just ignoring **OpenACC**)
 - Debugger/profiler use

MACC: Communication

- We leverage upper/lower-bounds of write/read accesses → *USE/DEF section*
 - Based on affine property, the sections are dynamically calculated for each combo of {GPU, Kernel, Array} by using upper/lower-bounds of loop-counters (MACC embeds runtime components)
 - Non-affine accesses (e.g. $A[B[s]] = 1;$) indicate the entire array
- MACC's routines manage updated regions for each combo of {GPU, Array} → *DIRTY section*
 - Before a kernel execution, necessary communications are deduced by the superposition of sections
 - Communications are executed through host memory removing duplicated transfers, or via NVLink

MACC: Communication



- Multi-GPU execution is enabled when the kernel's all DEF sections don't overlap among GPUs
 - The switch between single/multi-GPU execution is performed at runtime involving communications

MACC's Extension: Polyhedral Compilation

- To complement our analysis, we employ PLUTO to perform loop-fission before transpilation

```
#pragma acc parallel loop
for (j1 = 0; j1 < M; j1++)
#pragma acc loop
  for (j2 = j1; j2 < M; j2++) {
    symmat[j1][j2] = 0.0;
    for (i = 0; i < N; i++)
      sysmat[j1][j2] +=
        data[i][i1] * data[i][i2];
    sysmat[j2][j1] = sysmat[j1][j2];
  }
```

```
#pragma acc kernels
{
  for (j1 = 0; j1 < M; j1++)
    for (j2 = j1; j2 < M; j2++) { /*...*/ }
  for (j1 = 0; j1 < M; j1++)
    for (j2 = j1; j2 < M; j2++) { /*...*/ }
  for (j1 = 0; j1 < M; j1++)
    for (j2 = j1; j2 < M; j2++) { /*...*/ }
}
```

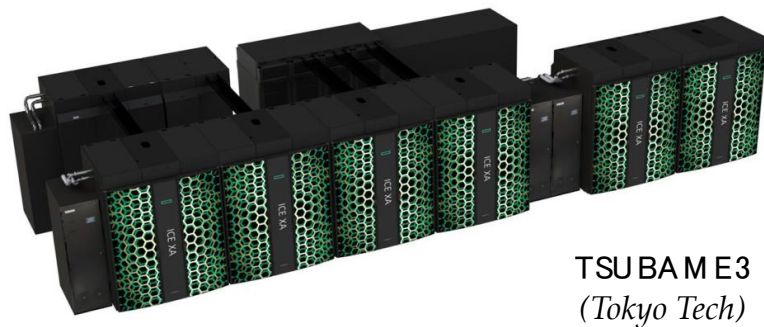
Try

PLUTO (--no-fuse)

Reshape & Add #kernels

Implementation & Evaluation

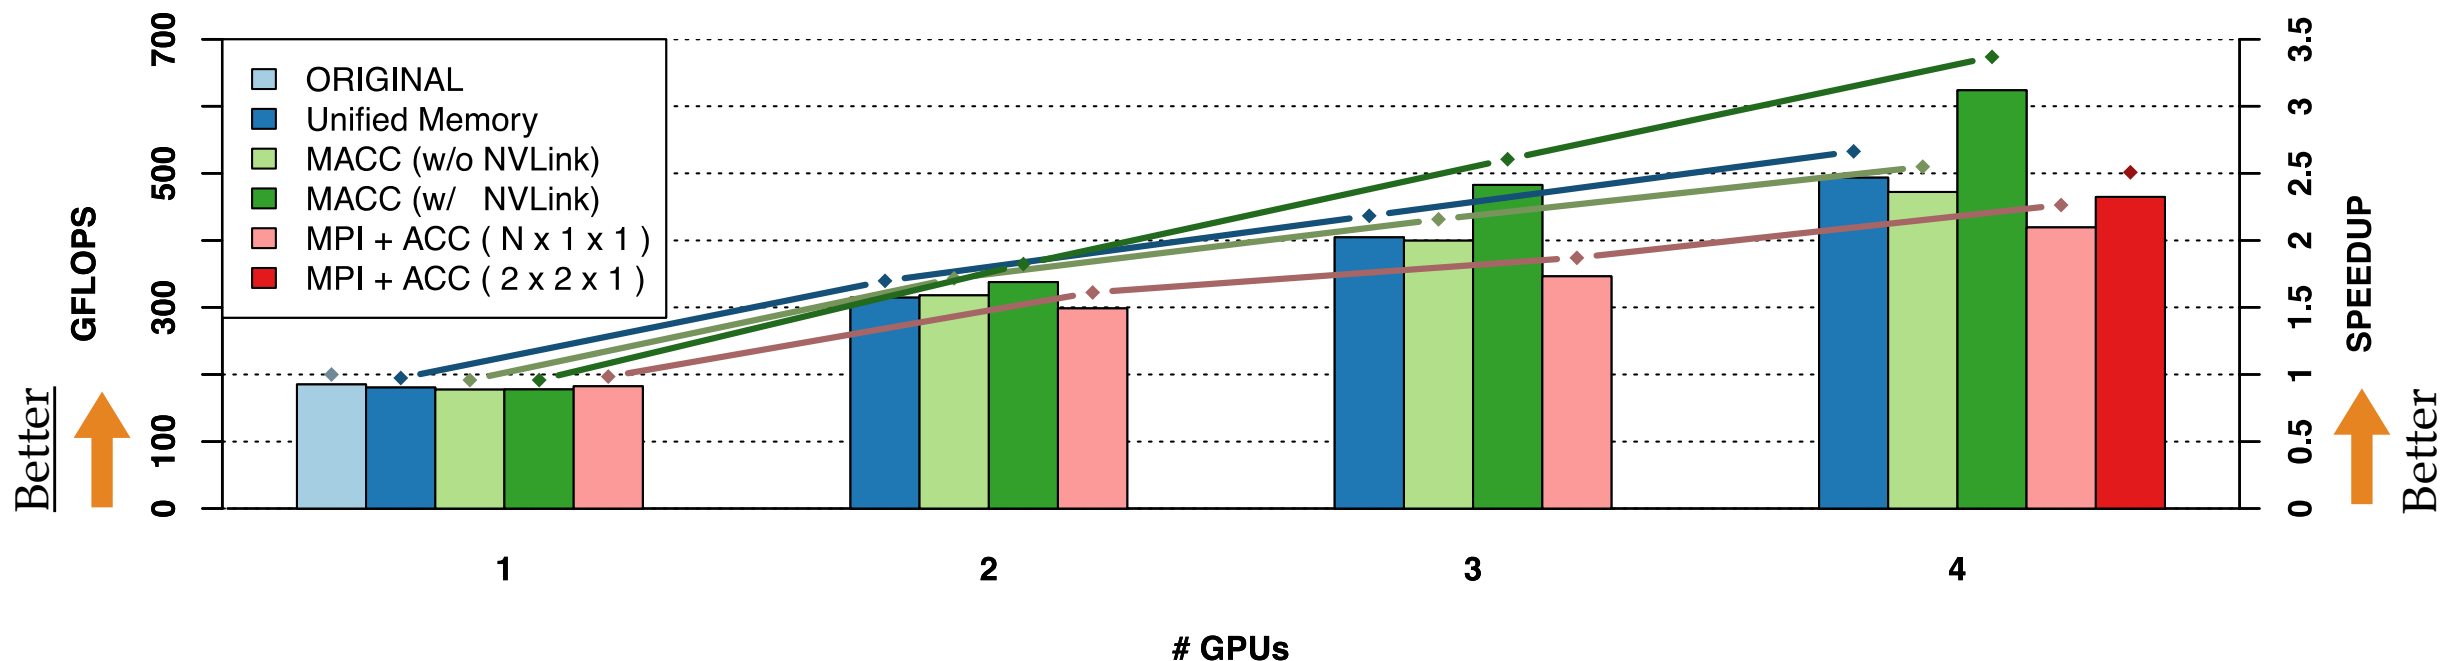
- We implemented MACC's prototype as a converter of XcodeML/C (which represents C code in XML)
 - Currently multi-dimensional arrays are treated as 1-dimentional arrays
- We measured performances of several benchmarks written with OpenACC on one node of TSUBAME3
 - Comparing to MPI+OpenACC version and Unified Memory version



CPU	Intel Xeon E5-2680 V4 (Broadwell-EP 14core) x 2
GPU	NVIDIA P100 (16GB HBM2@732GB/s) x 4
Compiler	PGI Compiler 17.10
CUDA	CUDA 9.0
NVLink	GPU0 ⇔ GPU2, GPU1 ⇔ GPU3: 40GB/s (one-way) Others: 20GB/s (one-way)

Himeno (Iterative 19-point Stencil Computation)

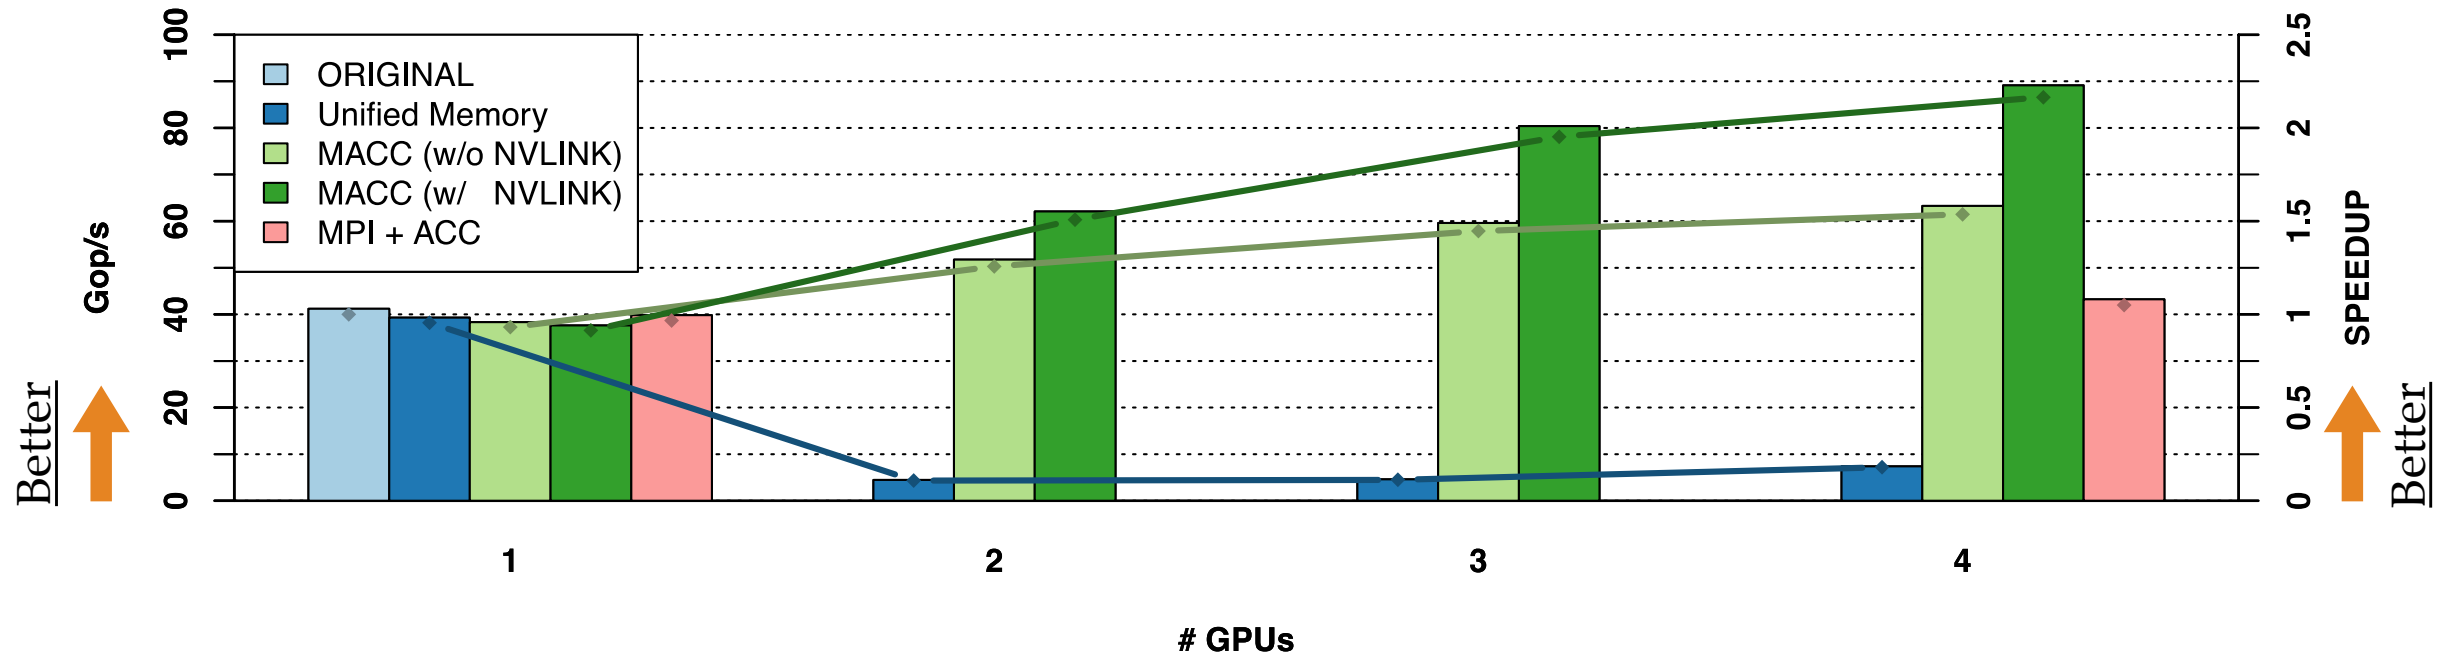
Size: $(i, j, k) = (256 \times 256 \times 512)$, Halo Communication (approx. $255 \times 511 \times 8$ bytes)



- MACC (w/ NVLink) achieved 3.36× speedup (32.1% performance increase compared to no NVLink)
- Unified Memory, that uses NVLINK, is slightly better than MACC (w/o NVLink)

NPB-CG (Iterative SpMV + Eigenvalue Calculation)

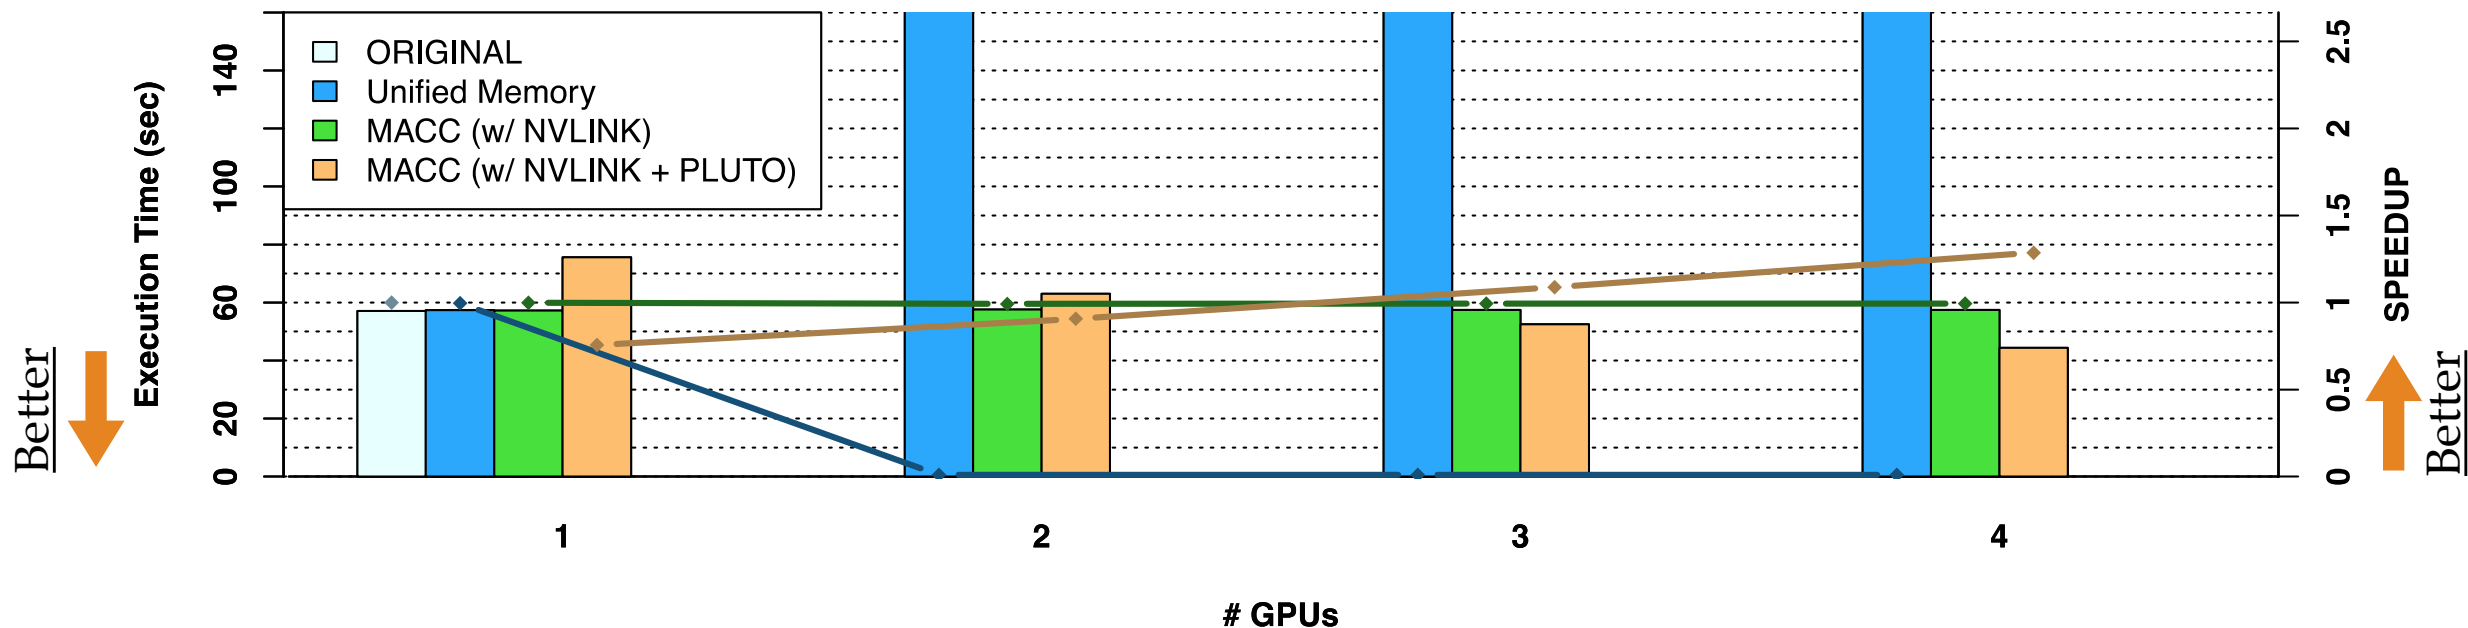
Size: rowsize = 150,000, All-to-All Communication (rowsize / GPUNUM × 8 bytes)



- MACC (w/ NVLink) gained the highest performance (40.9% performance increase compared to no NVLink)
- Unified Memory degraded the performance due to memory thrashing (frequent page fault & migration)
- MPI version (limited to $\text{proc}=n^2$) had low performances due to redundant communications

PolyBench/ACC COVAR (Covariance Matrix Calculation)

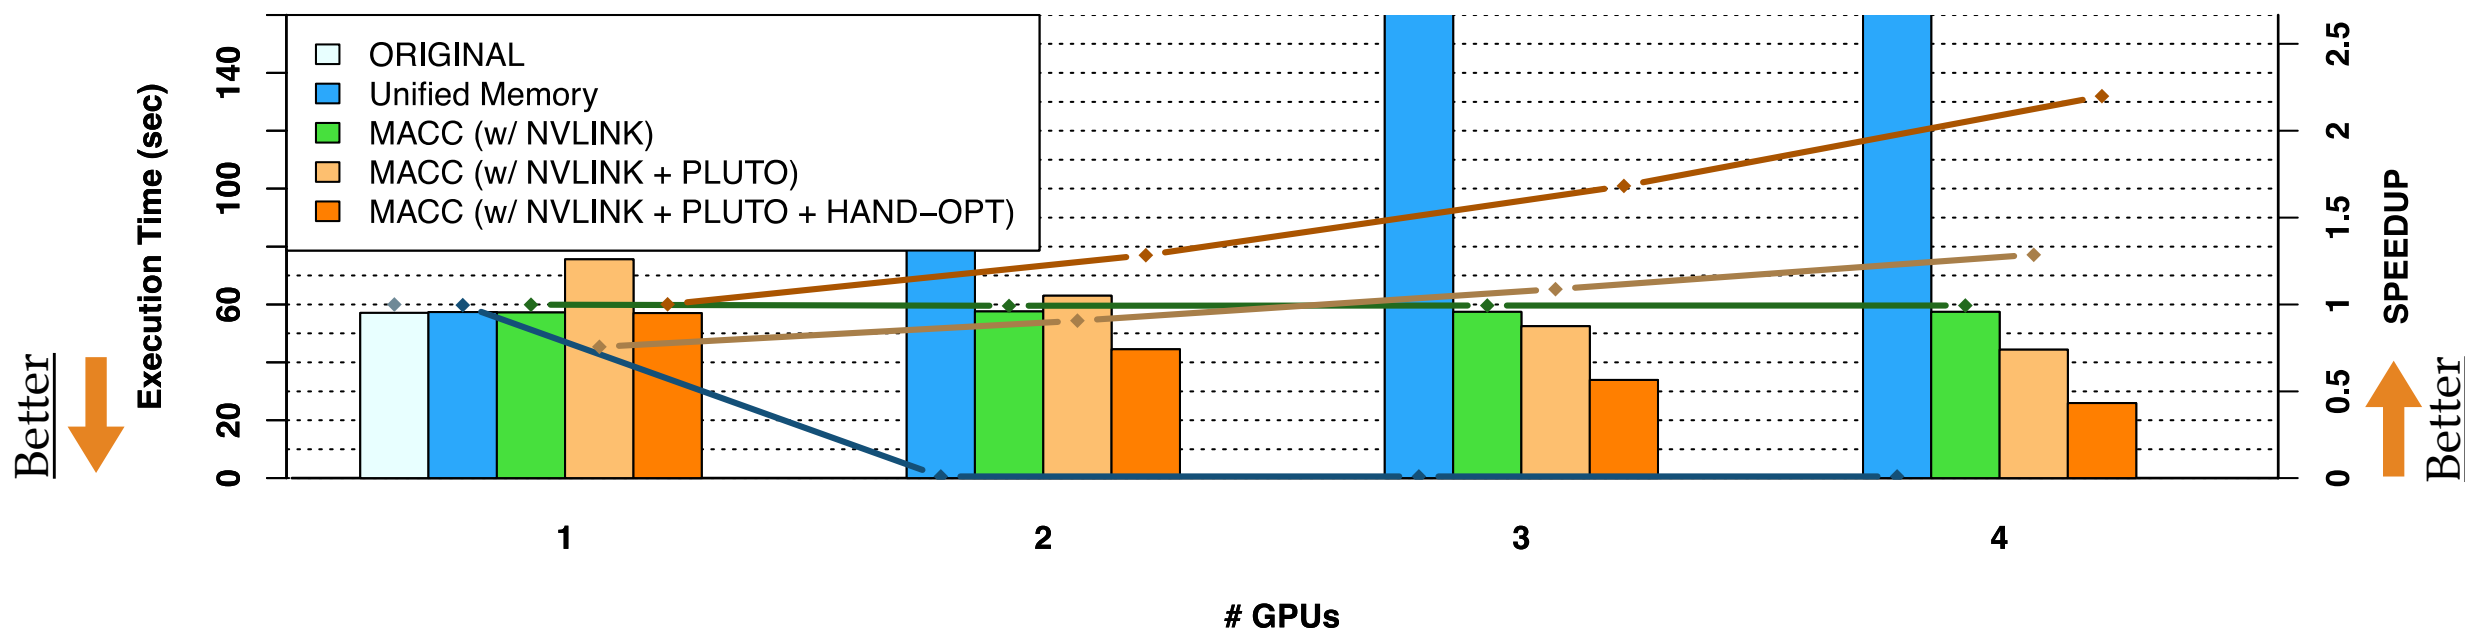
Size: 16,348 × 16,348



- A kernel combining symmetric-matrix creation (SC) + covariance calculation (CC) prevented multi-GPU execution
- After loop-fission by PLUTO, the SC and the CC were separated (1.29× speedup when using four GPUs)
 - The CC was executed on multi-GPU
 - The SC was executed **sequentially** ☹, because of MACC and PGI's poor loop-carried dependency checker

PolyBench/ACC COVAR (Covariance Matrix Calculation)

Size: 16,348 × 16,348



- The SC was executed **sequentially** ☹, because of MACC and PGI's poor loop-carried dependency checker
 - To solve the dependency, we added one directive line of loop construct into source-code after PLUTO
 - Still the SC was executed on single-GPU, but we achieved 2.20× speedup
 - We will need a loop model even in loop-dependency checker

Conclusion

- We built an OpenACC transpiler to use multi-GPU automatically
 - keeping the semantic and the portability
 - Communications are generated based on upper/lower-bounds of array accesses
- 3.36× speedup with stencil, and 2.16× speedup with NPB-CG when using four GPUs
- GPU-to-GPU communication via NVLink improved the performances
- Future work:
 - More analysis (affine and non-affine program analysis)
 - Work-sharing optimization (temporality, fine distribution)
 - Combining with task-based system
 - More accelerators, Hetero computing



Tokyo Tech

Thank you





Tokyo Tech

Back up



OpenACC (Execution Model)

- **#pragma acc parallel** : specifies a region executed on the accelerator (parallel region)

```
#pragma acc parallel
{ /* parallel region */ }
```

Appendable clauses:

```
num_gangs, num_workers, vector_length, if, async, reduction, ...
```

- **#pragma acc loop** : describes the parallelism of a loop
(automated, but still necessary for the performance improvement)

```
#pragma acc parallel num_gangs(8) num_workers(128) vector_length(128)
#pragma acc loop independent gang reduction (+ : sum)
for (int i = 0; i < N; i++)
#pragma acc loop independent worker
    for (int j = 0; j < N; j++)
#pragma acc loop independent vector
        for (int k = 0; k < N; k++) { /* ... */ }
```

Appendable clauses:

```
gang, worker, vector,
independent, seq, collapse,
reduction, ...
```

architecture-independent parallelism

Note: OpenACC provides abbreviated notations mixing several directives

OpenACC (Execution Model)

- **#pragma acc kernels** : treats loops as kernels, as far as possible
 - PGI Compiler mainly treats a tightly nested loop as a kernel

```
#pragma acc kernels
{
    for (int t = 0; t < MT; t++)
        for (int i = 0; i < MI; i++)
            a[i] += b[t] * a[i];
}
```

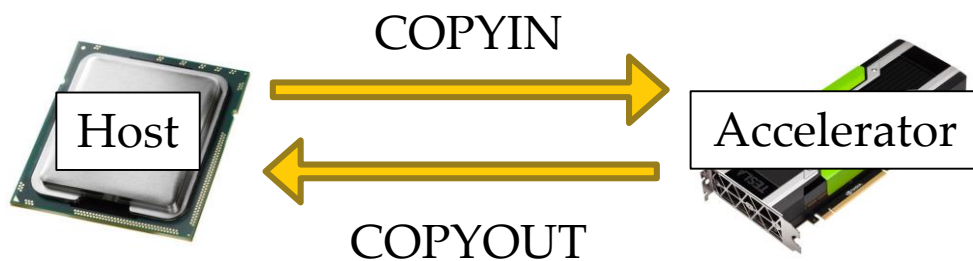
OpenACC Kernel: a nested loop executed in parallel on the accelerator

Tightly Nested Loop: nested loop which has just one statement inside except for the innermost

OpenACC (Memory Model)

- **#pragma acc data** : defines variables on the accelerator
 - A GPU manipulates its own memory (called **device memory**), so dependent data must be transferred
 - Before and after the region, communications corresponding specified clauses are occurred

```
#pragma acc data create(a[X:L]) copyin(b[X:L]) copyout(c[X:L]) copy(d[X:L]) present(e)
{
    /* Array 'a' 'b' 'c' 'd' and 'e' live here */
}
```



CREATE: Allocates memory space without transfers

COPY: COPYIN + COPYOUT

PRESENT: Already exists

- **#pragma acc update** : updates variables being defined on the accelerator (COPYIN or COPYOUT)

```
#pragma acc update host(a[start:length]) device(b[start:length])
```

MACC: Actual Example

Before

```
#pragma acc parallel loop gang reduction (+ : sum) present(a)
for (i = X; i < Y; i++) {
    sum += a[i + p];
}
```

PARALLEL REGION

After

```
{
    static int sections_are_changed = 1;
    sections_are_changed =
        (sections_are_changed || last_p != p || last_X != X || last_Y != Y);

    if (sections_are_changed) {
        section_are_changed = 0; last_p = p; last_X = X; last_Y = Y;

        calc_loop_sections(loop_sections, X, Y,
                           1 /* increment */,
                           0 /* whether to execute when X==Y */);

        init_uses(a_uses, 1 /* affine */); init_defs(a_defs, 0 /* none */);
        for (i = 0; i < NUMGPUS; i++) {
            update_section(a_uses[i], loop_sections[i].lb + p);
            update_section(a_uses[i], loop_sections[i].ub + p);
        }

        if (is_overlapping(a_defs)) {
            /* reconstruct for single GPU execution */
        }
    }
}
```

SECTION CALCULATION

```
#pragma omp parallel num_threads(NUMGPUS) reduction (+ : sum) private (i)
{
    int tnum = omp_get_thread_num();
    set_gpu_num(tnum);
```

```
    set_data_section(tnum, a, a_uses, a_defs);
    #pragma omp barrier
```

COMMUNICATION

```
#pragma acc parallel present(a)
#pragma acc loop gang reduction (+ : sum)
```

```
for (i = loop_sections[tnum].lb; i <= loop_sections[tnum].ub; i++) {
    sum += a[i + p];
}
```

PARALLEL REGION



Tokyo Tech